


Artificial Intelligence

Chapter 2


Agents



Updates and Additions: Dr. Siamak Sarmady

By: Dan Klein and Pieter Abbeel
University of California, Berkeley



Designing Agents





Demo: HISTORY – MT1950.wmv

Designing Rational Agents

- An **agent** is anything that can be viewed as **perceiving** its **environment** through **sensors** and **acting** upon that environment through **actuators**.
- In short: An **agent** is an entity that *perceives and acts*.
- A **rational agent** selects actions that maximize its (expected) **utility**.
- Characteristics of the **percepts**, **environment**, and **action space** dictate techniques for selecting rational actions and also the complexity of an agent (thermostat, a complex robot)



Pac-Man as an Agent




Pac-Man as an Agent



Example Agents


- **Human agent:** eyes, ears, and other organs for sensors; hands, legs, mouth, and other body parts for actuators
- **Robotic agent:** cameras and infrared range finders for sensors; various motors for actuators

Agents and environments



- The **agent function** maps from **percept histories** to actions:
$$[f: \mathcal{P}^* \rightarrow \mathcal{A}]$$
- If we can specify the agent's choice of action for every possible percept sequence, then we have said more or less **everything there is to say** about the agent.

Vacuum-cleaner world



- Percepts: location and contents, e.g., [A,Dirty]
- Actions: *Left, Right, Suck, NoOp*

A vacuum-cleaner agent

Percept sequence	Action
$[A, Clean]$	<i>Right</i>
$[A, Dirty]$	<i>Suck</i>
$[B, Clean]$	<i>Left</i>
$[B, Dirty]$	<i>Suck</i>
$[A, Clean], [A, Clean]$	<i>Right</i>
$[A, Clean], [A, Dirty]$	<i>Suck</i>
\vdots	\vdots

```
function REFLEX-VACUUM-AGENT([location,status]) returns an action
  if status = Dirty then return Suck
  else if location = A then return Right
  else if location = B then return Left
```

What is the right function?
Can it be implemented in a small agent program?

Agents Architecture

- The **agent function** is an **abstract mathematical description**; the agent program is a concrete implementation, running on the agent architecture.

Agent = program + architecture

- The **agent program** runs on the physical **architecture** to produce f

Architecture = sensors + actuators

Rational agents

- An agent should strive to "do the right thing", based on what it can perceive and the actions it can perform. The right action is the one that will cause the agent to be most successful
- **Performance measure:** An objective criterion for success of an agent's behavior.
- E.g., performance measure of a vacuum-cleaner agent could be: amount of dirt cleaned up, amount of time taken, amount of electricity consumed, amount of noise generated, etc.
- **Aim:** A rational agent wants to maximize the performance measure.

Rational agents

Assume an agent has committed an action. We want to assess whether it has been rational?

- **Rationality depends on**
 1. The performance measure that defines the criterion of success
 2. The agent's prior knowledge of the environment (i.e. sensor capabilities)
 3. The actions that agent can perform (i.e. actuator capabilities)
 4. The agent's percept sequence to date
- **Rational Agent:** For each possible percept sequence, a rational agent should select an action that is expected to maximize its performance measure, given the evidence provided by the percept sequence and whatever built-in knowledge the agent has.

Rational agents

- **Omniscience vs. rational:** Rationality is distinct from omniscience (all-knowing with infinite knowledge). An omniscience agent knows the actual outcome of its actions and can act accordingly; but omniscience is impossible in reality.*
- **Information gathering and exploration:** Agents can perform actions in order to modify future percepts so as to obtain useful information.
- **Autonomous agent:** An agent is autonomous if its behavior is determined by its own experience (with ability to learn and adapt) rather than the knowledge of its designers

PEAS

- Before designing an agent, we must first specify the setting for the intelligent agent. These settings are sometimes called PEAS.
- PEAS:
 1. Performance measure (in addition to measuring success, it is used to improve the performance)
 2. Environment
 3. Actuators
 4. Sensors

PEAS

- Consider, e.g., the task of designing an automated taxi driver:
 - **Performance measure:** Safe, fast, legal, comfortable trip, maximize profits
 - **Environment:** Roads, other traffic, pedestrians, customers
 - **Actuators:** Steering wheel, accelerator, brake, signal, horn
 - **Sensors:** Cameras, sonar, speedometer, GPS, odometer, engine sensors, keyboard

PEAS

- Agent: Medical diagnosis system
 - **Performance measure:** Increase in patient health, costs, lawsuits
 - **Environment:** Patient, hospital, staff
 - **Actuators:** Screen display (questions, tests, diagnoses, treatments, referrals)
 - **Sensors:** Keyboard (entry of symptoms, findings, patient's answers)





PEAS

- **Agent:** Part-picking robot
 - **Performance measure:** Percentage of parts in correct bins
 - **Environment:** Conveyor belt with parts, bins
 - **Actuators:** Jointed arm and hand
 - **Sensors:** Camera, joint angle sensors





PEAS

- **Agent:** Interactive English tutor
 - **Performance measure:** Maximize student's score on test
 - **Environment:** Set of students
 - **Actuators:** Screen display (exercises, suggestions, corrections)
 - **Sensors:** Keyboard

Environment types

-  **Fully observable** (vs. partially observable): An agent's sensors give it access to the **complete** state of the environment at **each point** in time.
-   **Deterministic** (vs. stochastic): The next state of the environment is completely **determined** by the **current state** and the **action** executed by the agent. (If the environment is deterministic except for the **actions of other** agents, then the environment is **strategic**)
-  **Episodic** (vs. sequential): The agent's experience is divided into **atomic "episodes"** (each episode consists of the agent perceiving and then performing a single action), and the choice of action in each episode depends only on the episode itself (and won't effect decisions of next episode).

Environment types

-  **Static** (vs. dynamic): The environment is unchanged while an agent is deliberating. (The environment is **semidynamic** if the environment itself does not change with the passage of time but the agent's performance score does)
 -  **Discrete** (vs. continuous): A limited number of distinct, clearly defined **percepts and actions**.
-  **Single agent** (vs. multi-agent): An agent operating by itself in an environment.
-  **known**(vs unknown): the outcomes or outcomes probabilities for all actions are given.

Environment types

	Chess with a clock	Chess without a clock	Taxi driving
Fully observable	Yes	Yes	No
Deterministic	Strategic	Strategic	No
Episodic	No	No	No
Static	Semi	Yes	No
Discrete	Yes	Yes	No
Single agent	No	No	No

- The environment type largely **determines** the agent design
- The **real world** is (of course) partially observable, stochastic, sequential, dynamic, continuous, multi-agent

Agent functions and programs

- An agent is **completely specified** by the agent function i.e. mapping percept sequences to actions
- An agent function (or a small equivalence class) is rational
- **Aim:** find a way to implement the rational agent function concisely
- **Key challenge** for AI is to find, how to **write program** that produce rational behavior from a **small** amount of code rather than from a large number of table entries.

Table-lookup agent

```
function TABLE-DRIVEN-AGENT(percept) returns action  
  static: percepts, a sequence, initially empty  
           table, a table, indexed by percept sequences, initially fully specified  
  
  append percept to the end of percepts  
  action  $\leftarrow$  LOOKUP(percepts, table)  
  return action
```

Figure 2.5 An agent based on a prespecified lookup table. It keeps track of the percept sequence and just looks up the best action.

- **Drawbacks:**
 - Huge table ($\Sigma^T P^T$ entries)
 - Take a long time to build the table
 - No autonomy
 - Even with learning, need a long time to learn the table entries

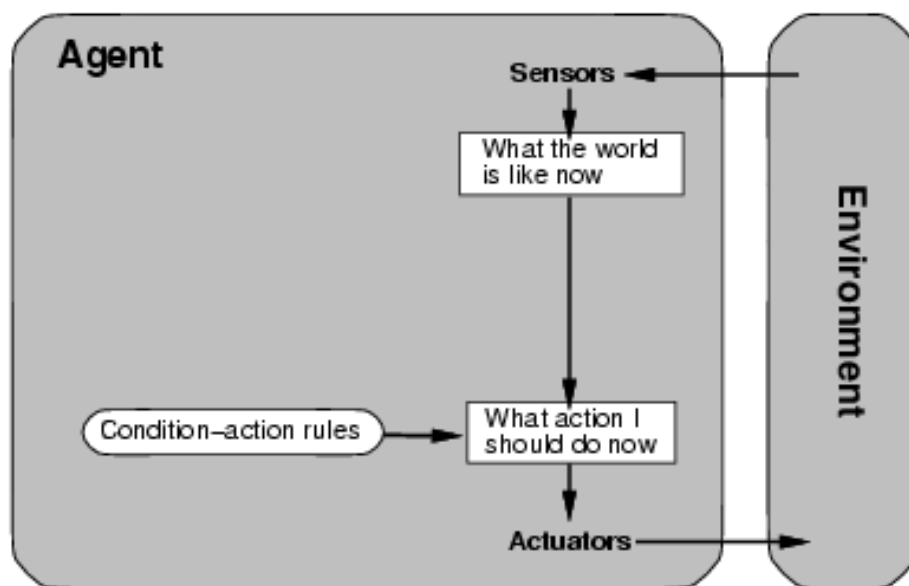
Agent program for a vacuum-cleaner agent

```
function Reflex-Vacuum-Agent( [location,status]) returns an action  
  if status = Dirty then return Suck else if location = A then return Right else if  
    location = B then return Left
```

Agent types

- Four basic types in order of increasing generality:
 - Simple reflex agents
 - Model-based reflex agents
 - Goal-based agents
 - Utility-based agents
- Learning Agents

Simple reflex agents



Simple reflex agents

```
function SIMPLE-REFLEX-AGENT(percept) returns action  
static: rules, a set of condition-action rules  
  
state  $\leftarrow$  INTERPRET-INPUT(percept)  
rule  $\leftarrow$  RULE-MATCH(state, rules)  
action  $\leftarrow$  RULE-ACTION[rule]  
return action
```

Figure 2.8 A simple reflex agent. It works by finding a rule whose condition matches the current situation (as defined by the percept) and then doing the action associated with that rule.

- These agents select actions on the basis of only the current percept, *ignoring* the rest of the percept history (e.g. vacuum agent, because based on current location and whether there it contains dirt).



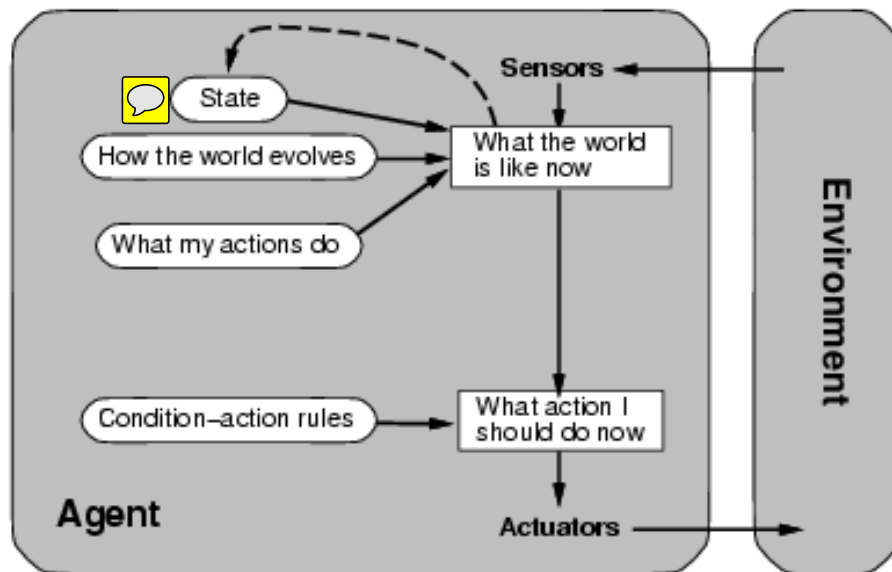
Simple reflex agents

- Vacuum agent has just 4 percepts in this way (down from 4^T).
- Condition-action rules (i.e. the first rule in the set that matches)

If car-in-front-is-braking then initiate braking

- Rectangles denote current internal state of the agent's decision process, the ovals represent the background information used in the process.
- Works only if the correct decision can be made on the basis of only current percept, that is only if the environment is fully observable.

Model-based reflex agents



Model-based reflex agents

- Keeps track of the part of the world it can see now i.e. maintains **some sort of internal state** that depends on the **percept history** (even though does not keep a percept history)
- Updating state information requires two kind of knowledge:
 - **First:** some info on **how** the world **evolves independently** of the agent, e.g. where will be the car nearby given we knew its location a while back
 - **Second:** **how** the **agent's** own actions **affect the world**, e.g. if the agent turns the steering wheel toward north, we will be 5 miles north in 5 minutes

Model-based reflex agents

function REFLEX-AGENT-WITH-STATE(*percept*) **returns** *action*

static: *state*, a description of the current world state
rules, a set of condition-action rules

state ← UPDATE-STATE(*state*, *percept*) (state, action, percept)

rule ← RULE-MATCH(*state*, *rules*)

action ← RULE-ACTION[*rule*]

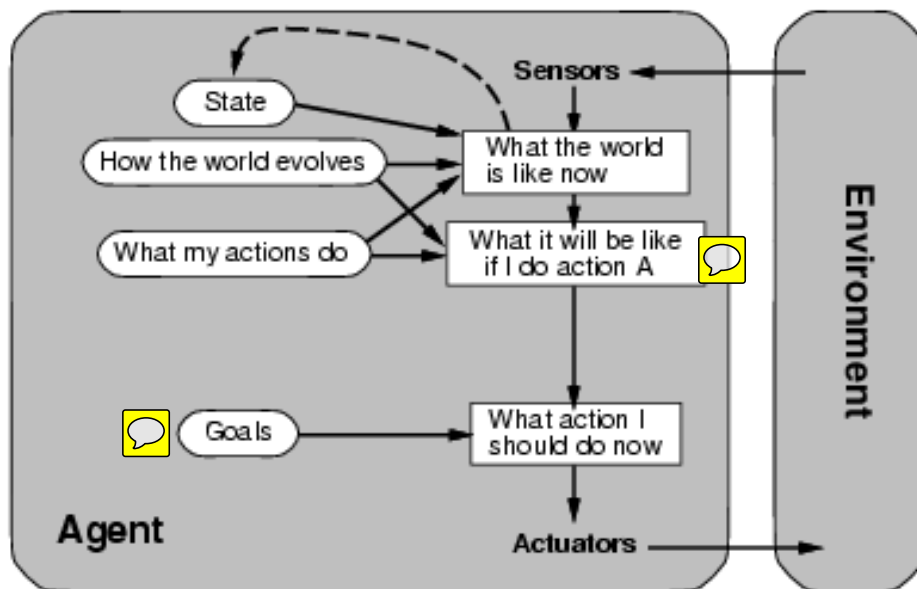
state ← UPDATE-STATE(*state*, *action*)

return *action*

Figure 2.10 A reflex agent with internal state. It works by finding a rule whose condition matches the current situation (as defined by the percept and the stored internal state) and then doing the action associated with that rule.

- Current **percept** is **combined** with the **old state** to generate the updated description of the **current state**.
- Update-State takes into account the previous action (In fact, the **new percept**, gives insight about the **effect of agent's action** on the state of environment).

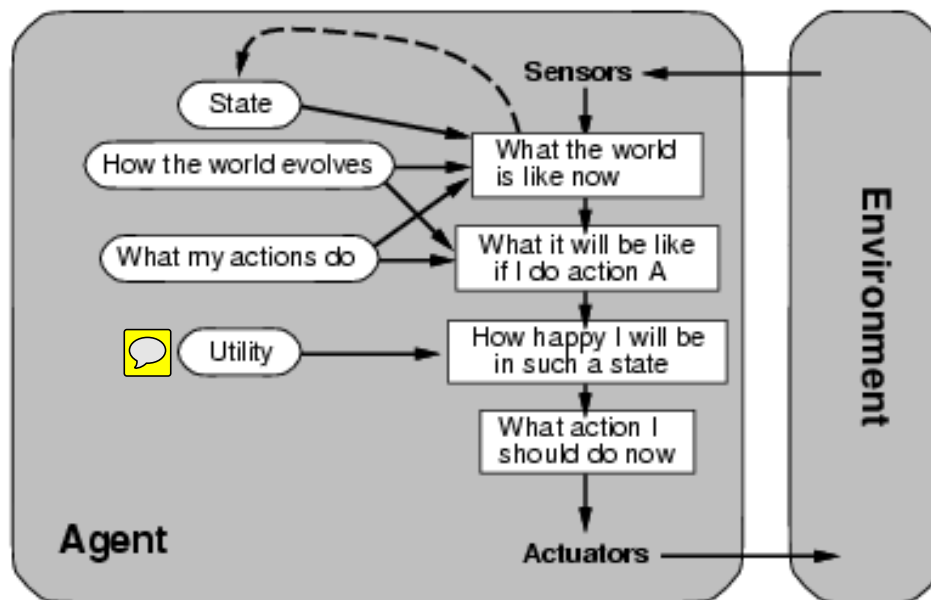
Goal-based agents



Goal Based Agent

- In addition to a current state description, the agent needs some sort of **goal information** that describes **situations that are desirable**. (e.g. for Taxi, the correct decision depends on where the taxi wants to go).
- Search and planning in the AI are devoted to find actions that achieve the agent's goals.
- Decision making of this kind is **fundamentally different** from rule based method. It involves the future, both "**what will happen** if I do such and such" and "**Will that make me happy**"?

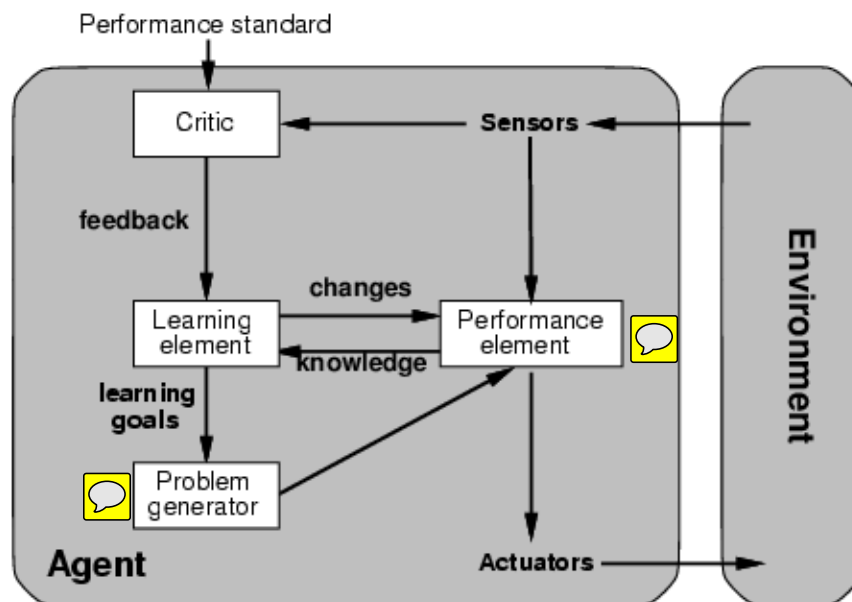
Utility-based agents



Utility Based Agent

- There are multiple action sequences that will get a taxi to destination (i.e. achieving goal) but some are quicker, safer, more reliable, or cheaper
- Goals just provide a **crude binary distinction** between "happy" and "unhappy" state. **Utility function** is a more general performance measure which produces a **real number** representing associated level of happiness.
- If one world state is preferred to another, then it has a higher utility for the agent.

Learning agents



Learning Agent

- In most areas of AI, it is the preferred method to build **learning machines** and then **teach them**.
- It allows the agent to operate in **initially** unknown environments and to **become more competent** than its **initial knowledge** alone might allow.
- **Performance (Execution) element:** takes the percepts and decides the action (the entire agent of previous types).
- **Learning element:** responsible for making improvements. Uses feedback from **Critic** to identify how the performance element should be modified to do better in the future. Its design depend very much on the design of performance element.

Learning Agent

- **Critic:** is necessary because the percepts themselves do not provide indication of agent's success. It is necessary that performance standard be **fixed** (outside the control of agent, so that agent cannot modify it to fit its behavior).
- **Problem Generator:** Responsible for suggesting **exploratory actions** that will lead to new and informative experiences.
- **Learning element:** can formulate a rule saying this was a bad action and it then installs the new rule in the performance element (modifies it).

Course Topics

- Part I: Making Decisions

- Uninformed Search
- Informed Search
- Constraint satisfaction
- Adversarial search

- Part II: Logic

- Logical Agents
- 1st Order Logic
- Resolution using 1st Order Logic

- Additional

- Machine Learning

