


Artificial Intelligence


Chapter 6: Adversarial Search



Updates and Additions: Dr. Siamak Sarmady
By: Dan Klein and Pieter Abbeel
University of California, Berkeley

Game Playing

- If **someone** or a computer can **play** a game with you and can **win**, it shows **smartness** and intelligence.
- So researchers have used games to **make** computers **smart** and show that they are **intelligent**...



Checkers:

1950: First computer player.

1994: First computer champion, Chinook, ended 40-year dominance of human champion Marion Tinsley using complete 8-piece endgame.


It means, it can understand all possible scenarios until next 8 moves, if there is a **guaranteed** possibility of win, enforce a win.

2007: Checkers solved!

It means your program can understand all possible scenarios from the beginning of the game (e.g. if you are the beginner), whether it is possible to enforce a win or a tie or you will lose ...

In checkers, a **tie** will happen if both play **correct** (tic-tac-toe is same). If the opponent make a **mistake**, you can enforce a win

Game Playing




Chess:

1997: Deep Blue defeats human champion Gary Kasparov in a six-game match.

Deep Blue examined **200M positions** per second, used methods that searched up to **40 steps ahead**. Current programs are even better.

We **still** cannot check the **whole game** (whether there is a guaranteed way to enforce a win or tie)

Game Playing



Go:


Human champions are now **starting** to be challenged by machines, though the best humans **still beat** the best machines.

In go, branching factor, **$b > 300!$** (i.e. moves in each step). So Classic programs use **pattern** knowledge bases, but big recent advances use **Monte Carlo** (randomized) expansion methods.

That is in this game computer software do **not look everywhere**, they just search **some of** the search space...

These programs are not in champion level, but in expert level.

Game Playing State-of-the-Art



- Pacman?
We don't know...


Behavior from Computation

Try to **guess** what **kind of computations** happen in the software to produce these behaviors you see?

It is **NOT** bunch of IF...Then commands!

- **Each player** (Eater and Ghosts) are **searching** for options at each step

- **One player** plays at each step and then **the others** decide what to do to **increase** the **chance** of win...



[Demo: mystery pacman (L6D1)]

Video of Demo Mystery Pacman


Try to **guess** what **kind of computations** happen in the software to produce these behaviors you see?

It is **NOT** bunch of IF...Then commands!


- **Each player** (Eater and Ghosts) are **searching** for options at each step

- One player plays at each step and then the others decide what to do to **increase** the **chance** of win...

Wasn't perfect?



Adversarial Games




Two or **more** agents play, **each** trying to **maximize** its **own** utility...

Sometimes, **one** player's **win**, is the **other's** **loss**


Types of Games

- Many different kinds of games!
- Possibilities:
 - Deterministic vs. stochastic
 - One, two, or more players
 - Zero sum vs. cooperative
 - Perfect information (Chess), Partial Info (Cards)
- We want algorithms for calculating a **strategy (policy)** which suggests a move in each state...
 - Strategy means series of actions
 - Since it is **not single player**, we **don't know** what is going to **happen**. In **each** situation we want to know **what** the agent should **do**...




Deterministic Games

- Formalization:
 - **States:** S is the set of states in the game (start at s_0)
 - **Players:** $P=\{1...N\}$ (usually take turns)
 - **Actions:** set of A (may depend on player / state)
 - **Transition Function:** $S \times A \rightarrow S$
 - Maps an s_n state to a resulting state S_{n+1} because of an action
 - We used to call it a **successor function** in single player...
 - **Terminal Test:** $S \rightarrow \{t, f\}$
 - A test that in last state, will tell us **whether** the agent **won**
 - **Terminal Utilities:** $S \times P \rightarrow R$
 - Determines the **utility** for each player in the last state
 - In games we normally calculate the utility **at the end** of the game
- **Solution** for a player is a **policy**: $S \rightarrow A$



Zero-Sum Games

Each one seeks higher utility for itself, but they simplify the job for the other as they sort, not a full collaboration though




Zero-Sum Games

- Agents have **opposite utilities** (values on outcomes)
- Lets us think of a **single value** that one maximizes and the other minimizes
- Adversarial, **pure competition**

General Games

- Agents have **independent utilities** (values on outcomes)
- Cooperation, indifference, competition, and more are **all possible**
- More later on non-zero-sum games

Adversarial Search



Decisions are taken based on what each agent **thinks** the other agent **will do**...


So, assuming we perform the **first move**, we think **what actions** the other agent **might perform**, and in answer to each of those actions **what we should do** ...

In **some** games it is **possible** to have a **search tree**

Single-Agent Trees

One move: -1
 Eat food pallet: +10
 Sometimes game win: +500

To which side do you think the agent will go, if it has these utilities?




Value^(potential) of a State (Single Player)

Value of a state:
 The **best** achievable outcome (utility) from that state

We do **maximum propagation**, because the agent has the full decision and it will naturally want to select the node with higher utility...


So we build "maximum utility direction" information for it



Propagate the **best** outcomes towards up...

Terminal States:
 $V(s) = \text{known}$

Adversarial Game Trees



Here, if both players play carefully, it is **guaranteed loss** for Pacman

(track moves to see)

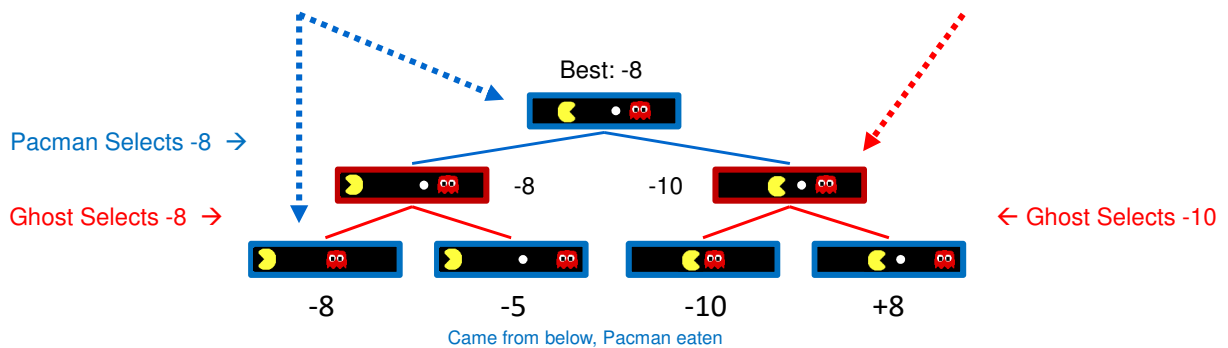
Minimax Values

Selection of Next States Under Agent's

Control: $V(s) = \max_{s' \in \text{successors}(s)} V(s')$

Selection of Next States Under Opponent's

Control: $V(s') = \min_{s \in \text{successors}(s')} V(s)$



Minimizing Node: Ghost selects the **next move**


Maximizing Node: Pacman selects the **next move**

Propagate up, for maximizing nodes select max, for minimizing nodes take min

Terminal States:


$V(s) = \text{known}$

Tic-Tac-Toe Game Tree



Adversarial Search (with Minimax) Summary

- Deterministic, zero-sum games:**
 - Tic-tac-toe, chess, checkers, Pacman
 - One of the players** maximizes result
selects the node with Max utility
 - The other** minimizes result
selects the node with Min utility of opponent, i.e. max of his own
- Minimax search algorithm:**
 - A state-space search tree
 - Players alternate turns
 - Compute every node's minimax value:**
 - Visit all branches and terminal states (DFS). Propagate the values upward and calculate the Minimax values.
 - Try to go a path with the best achievable utility against a rational (optimal) adversary



Max is not 8, ... You cannot force 8 in left branch to your opponent. Max that you can force is 5. You can still hope for 8 (if opponent makes mistake)

Minimax Implementation

```
def max-value(state):
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

Initialize with worse possible score...

And then use recursive calls to min-value and max-value in turns

Just works if we have **two agents** which play **in turns**

```
def min-value(state):
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Implementation (Dispatch Function)

Better implementation, decides based on the node type i.e. maximizer or minimizer...

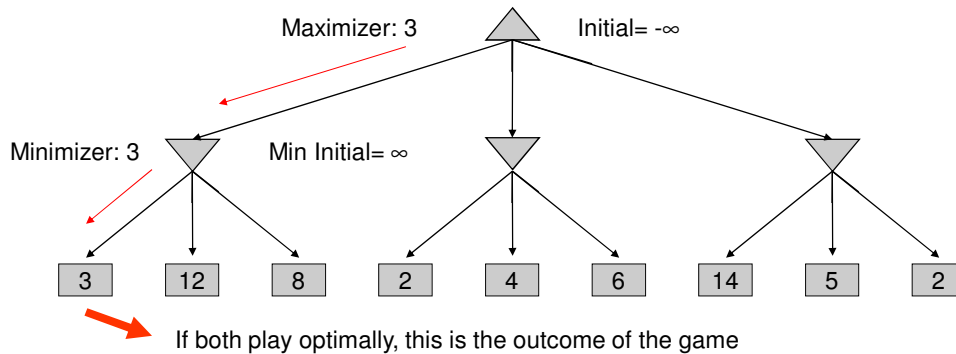
```
def value(state):
    if the state is a terminal state: return the state's utility
    if the next agent is MAX: return max-value(state)
    if the next agent is MIN: return min-value(state)
```

Because we don't always have a maximizer and a minimizer, e.g. if we have 3 players... i.e. pacman and 2 ghosts

```
def max-value(state):
    initialize v = -∞
    for each successor of state:
        v = max(v, value(successor))
    return v
```

```
def min-value(state):
    initialize v = +∞
    for each successor of state:
        v = min(v, value(successor))
    return v
```


Minimax Example



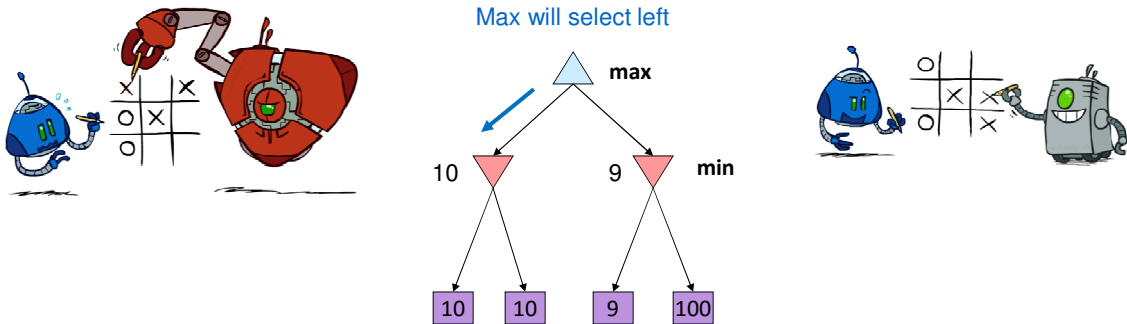
In **exam** you will **propagate** from bottom to top, but in programs you don't have/create the whole tree... you use DFS etc.

Minimax Efficiency

- How efficient is minimax?
 - Just like (exhaustive) DFS
 - Time:** $O(b^m)$, b : branching m : depth
 - Space:** $O(bm)$
- Example: For chess e.g. $b \approx 35$, $m \approx 100$
 - Exact** solution is completely **infeasible**
 - But, do we **need** to explore the **whole tree**?
Can we **get away** with **partial** search?



Minimax Properties



Optimal against a perfect player. Otherwise?

If we **suspect** that the minimizer may make **mistake**, perhaps we **go to right**? (9 and 10 are not that far, and we have chance of winning 100!)

[Demo: min vs exp (L6D2, L6D3)]

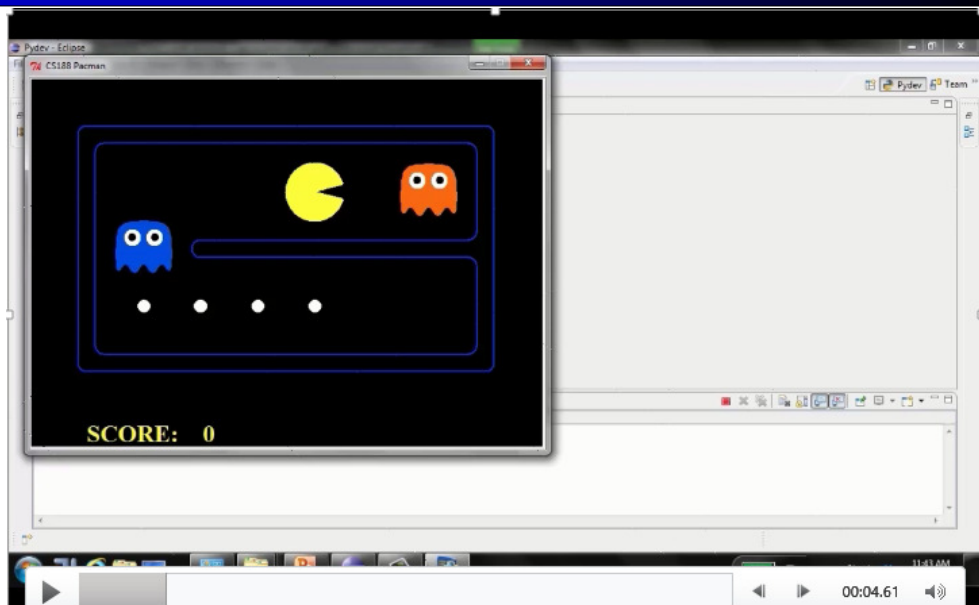
Video of Demo **vs. Expert** (Minimizer)

If both **ghosts** are **experts**, then we **already know** that there is no chance....

In that case **what should** packman **do** to maximize its utility?

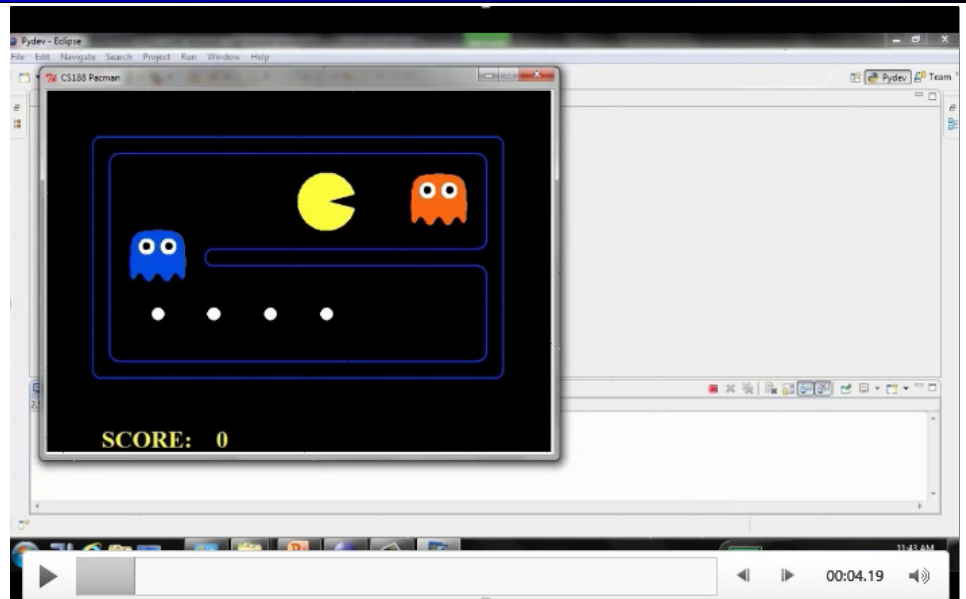
If **living** means **getting** more and more **negative** score, **better die** and stop the negative scores..!

Since we **lose 1** by **each move**, better suicide and **lose just 1** than wasting more scores...




Video of Demo vs. Non-Expert Ghosts

If the **ghosts** play **randomly**, then we have a 50% **chance** that the **blue** moves **down** (or up)... so we **take** our chance...




Resource Limits



Resource Limits

- **Problem:** In realistic games, **cannot search to leaves!**
- **Solution:** **Make up values!**
 - Instead, search only to a **limited depth** in the tree
 - Replace terminal utilities with an **evaluation function** for **non-terminal** positions...
 - We **wish** that those non-terminal **value represents the rest** of the branch (exactly like **heuristic**). But we cannot be sure of that...
 - **Quality** of the evaluation function is **important**
 - Now we just have **estimates**... **not real** MinMax values.
- **Example:**
 - Suppose we have **100 seconds**, can explore 10K nodes / sec
 - So can check **1M nodes** per move
 - With α - β pruning, can **reach** about **depth 8** – decent chess program
- **Guarantee of optimal play is gone**
- **Checking more levels makes a BIG difference in accuracy**
 - Goes near to actual end game scores (**otherwise** we would **only** evaluate **2 highest nodes**)
- **Use iterative deepening for real time algorithm** (until time runs up)



Depth Matters

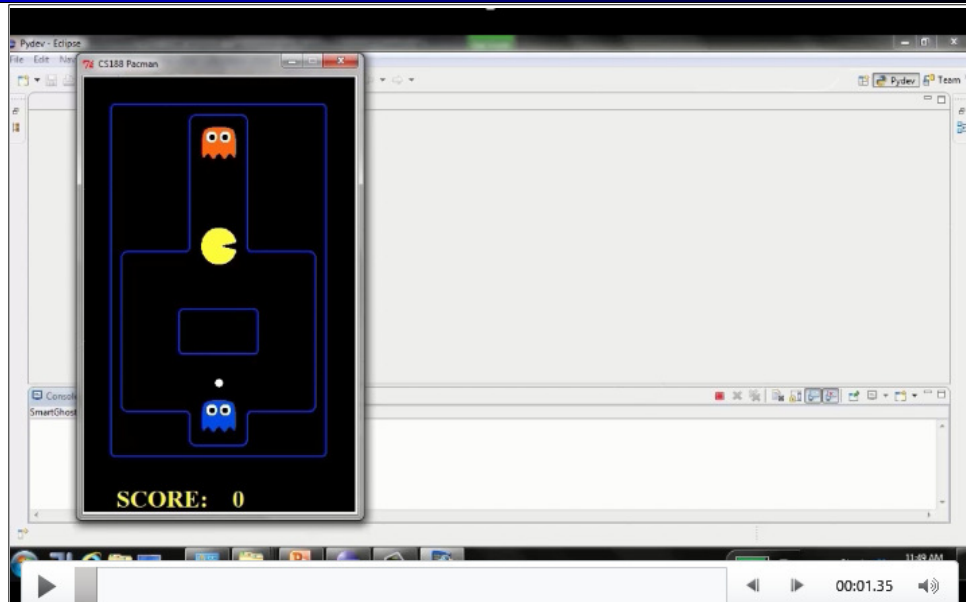
- Evaluation functions are always **imperfect**
- The **deeper** in the tree you go, your results will be more reliable and the **quality** of the evaluation function **matters less**



[Demo: depth limited (L6D4, L6D5)]

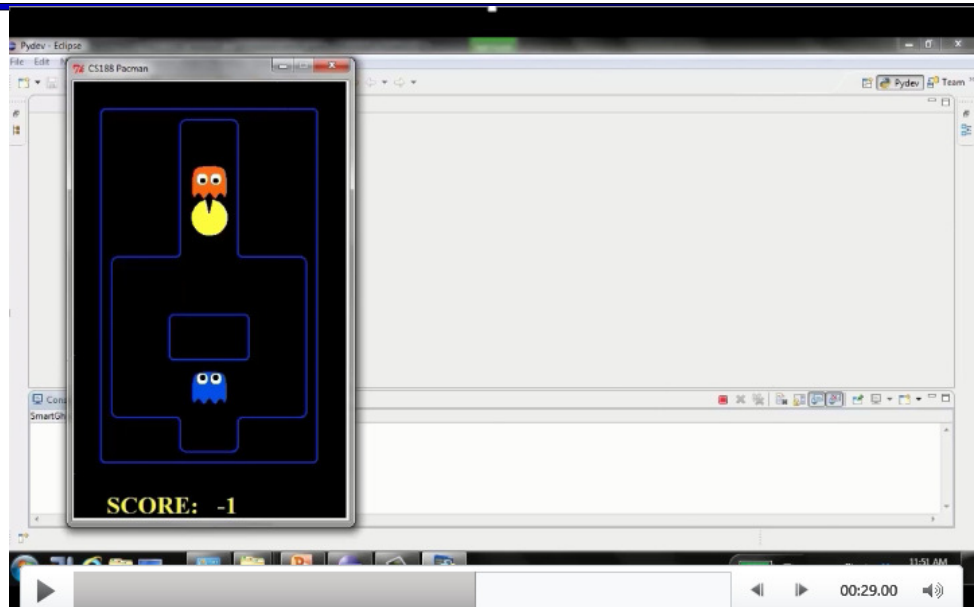
Video of Demo Limited Depth (2)

- In this demo we only search 2 levels deep
- Players cannot stop, they should move
- Pacman can either go up or down, if down it does not know which side will blue ghost go
- Pacman decides that it should not go up, but does not know which side to go after it comes down ...
- so blue closes the road to it, while orange closes its back




Video of Demo Limited Depth (10)


- This time we look ahead 10 moves
- This time pacman finds that the blue ghost should decide to go left or right after coming up
- It goes up one step to give time for itself to see what blue is going to do (it knows it can still be ahead of orange ghost)...
- Notice that the evaluation function is the same, only we search more steps ahead



Evaluation Functions



Evaluation for Pacman



- The evaluation function could be just a **running score** of the game ... so in Pacman it could be something like $f(x) = \text{Eaten Palets} * 10 - \text{Movements}$
- But **what** happens **if** our evaluation function is **not good** or **sufficient**
- It should be **able to differentiate** between states and gives **better** value for **better** states

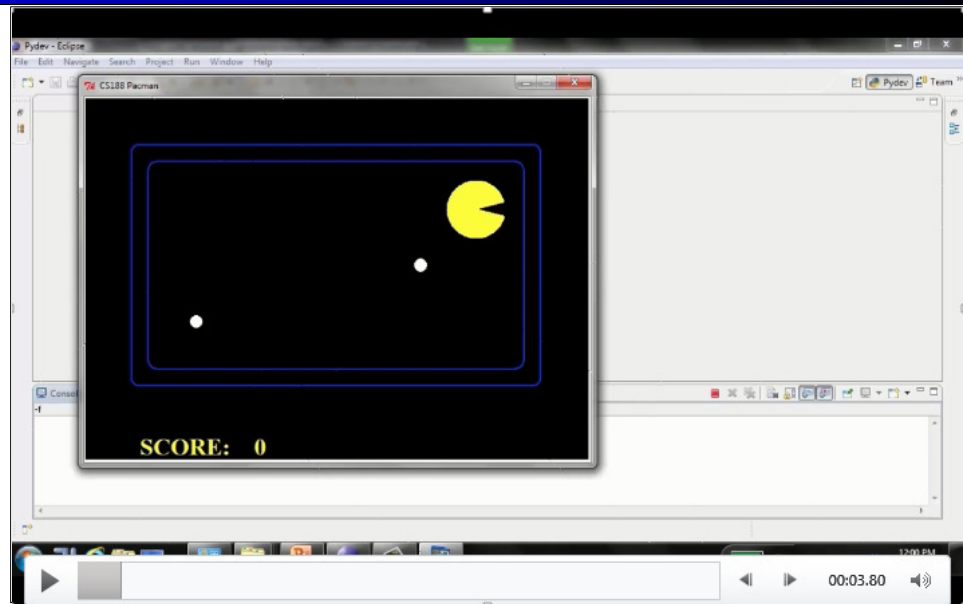
[Demo: thrashing d=2, thrashing d=2 (fixed evaluation function), smart ghosts coordinate (L6D6,7,8,10)]

Video of Demo Thrashing (d=2)

- Evaluation function is just **score** of game


- What happens if we just do **two levels** of evaluation only...

- there are no ghosts and you think it should be easy



Why Pacman Starves

- determine last layer's scores
- Then propagate upside
- depth=2




■ A danger of re-planning agents!

1. It knows his score will go up by eating the dot now (west, east)
2. It knows his score will go up just as much by eating the dot later (east, west)
3. There are no point-scoring opportunities after eating the dot (within the horizon, two here)
4. Therefore, waiting seems just as good as eating: he may go east, then back west in the next round of re-planning!

Video of Demo Thrashing -- Fixed (d=2)


So how we fix it?

- By using a better heuristic, for example **adding the distance** to food pallets to the function
- In that case the graph we draw will no more produce a tie...
- If you see **thrashing** in your problems, normally it is something in your **evaluation function**



Evaluation Functions

- Evaluation functions determine **scores** for **non-terminal** states in depth-limited search




- **Ideal function:** returns the **actual minimax** value of the position
- **In practice:** **approximate** it. For example, **sum up the scores** each wins (might represent how good a branch is)
- **Chess:** typically weighted linear sum of features (produces a real value)


$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- Such as: $f_1(s) = (\text{num white queens} - \text{num black queens})$, and f_2 for bishops etc.

Game Tree Pruning

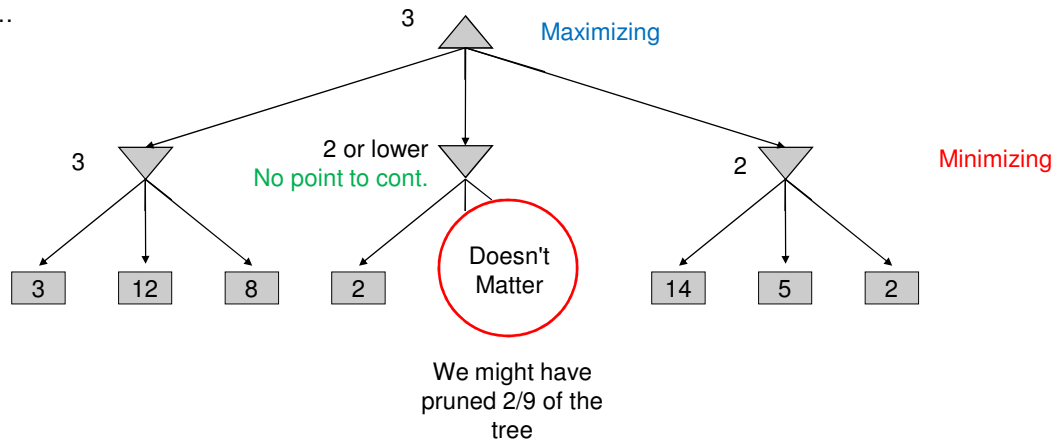


Minimax Example



Minimax Pruning

Maximizer will not choose anything lower than 3 so...




Alpha-Beta Pruning

General configuration (MIN version)

- We're computing the MIN-VALUE at some node n
- We're looping over n 's children
- n 's estimate of the childrens' min is dropping
- Who cares about n 's value? MAX
- Let a be the best value that MAX can get at any choice point along the current path from the root
- If $v \leq a$ and n becomes $\leq a$, MAX at the top will avoid reaching it, so we can stop considering n 's other children (it's already bad enough that it won't be played), so as soon as seeing $v=3$, (prune if $v \leq a$)

MAX version is symmetric

- The same way, if we are finding MAX-VALUE, the value cannot be $\geq b$ because the MIN in higher levels will avoid reaching the lower part of the tree (prune if $v \geq \beta$)



We keep best MAX (α) and MIN (β) value we have seen along the path up until the route

Alpha-Beta Implementation

α is the biggest value seen
 β is the smallest value seen

α : MAX's best option on path to root
 β : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):
    initialize  $v = -\infty$ 
    for each successor of state:
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$ 
        if  $v \geq \beta$  return  $v$ 
         $\alpha = \max(\alpha, v)$ 
    return  $v$ 
```


```
def min-value(state,  $\alpha$ ,  $\beta$ ):
    initialize  $v = +\infty$ 
    for each successor of state:
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$ 
        if  $v \leq \alpha$  return  $v$  //no need to cont.
         $\beta = \min(\beta, v)$ 
    return  $v$ 
```

Alpha-Beta Pruning Properties

- This pruning has **no effect** on minimax value computed for the root!
- Values of intermediate nodes might be wrong
 - **Important:** children of the root (**intermediate**) may have the **wrong value** **because** we perform **pruning** and do not always calculate exact values...
 - So the most naïve version won't let you do action selection
For example, will not help you in breaking ties for 10-10
 - Just passing up the alpha, beta will not be enough...


- Good child ordering **improves** effectiveness of pruning
- With "perfect ordering":
 - Time complexity drops to $O(b^{m/2})$... search better Childrens first?
 - **Halves** the depth so it can double the solvable depth!
 - Full search of, e.g. **chess**, is **still** hopeless...

- This is a simple example of **meta-reasoning** (computing about what to compute)



Alpha-Beta Quiz

Where does pruning happen?



Alpha-Beta Quiz 2

Where does pruning happen?

